End to End ML project

1. Get data :-

os.path.join → specifies the directory where the extracted data shd be downloaded and stored.

os.makedirs → makes directory.

* describe() and info() function to analyse data

* Median_house_value is capped.

* housing.hist(bins=50, figsize(20,15))
  - histogram is plotted which is used for univariate analysis & bins are used to specify the width of each group

* sklearn.train_test()
  It divides the dataset into train set & test set, test ratio can also be specified to divide the dataset

* train_set.shape
  gives shape of train set (columns & rows)

* housing['id'].value_counts()
  → frequency of id columns

2. Discover & Visualize the Data to gain insight
* strat_train_set.shape()
* We plot scatter plot by using
housing.plot(kind='scatter', x='longitude', y='latitude'). plt.show()

3. prepare the data for Machine learning
a) Data Cleaning
Get rid of whole attribute (feature)
housing.drop('total_bedrooms', axis=1)
b) Handling text & categorical Attributes
c) Custom Transformers
d) Feature scaling

4) Select & train a Model
Linear Regression, Decision Tree Regressor

5) Fine tune your model
Grid search CV is a method provided by Scikit-learn library in python for hyperparameter tuning of ML models.

6) Launch, Monitor & Maintain your system.
* automate by:-
- Collecting fresh data regularly & labelling it
- writing script to train model & fine tune the hyperparameters
- writing script to evaluate the model

13/11/24

# Python Implementation of Linear Regression

```
import numpy as np
import matplotlib.pyplot as plt

def estimate_coef(x, y):
    n = np.size(x)
    m_x = np.mean(x)
    m_y = np.mean(y)

    SS_xy = np.sum(y*x) - n*m_y*m_x
    SS_xx = np.sum(x*x) - n*m_x*m_x

    b_1 = SS_xy / SS_xx
    b_0 = m_y - b_1*m_x
    return (b_0, b_1)


def plot_regression_line(x, y, b):
    plt.scatter(x, y, color="m", marker="o", s=30)
    y_pred = b[0] + b[1]*x
    plt.plot(x, y_pred, color="g")
    plt.xlabel('x')
    plt.ylabel('y')

def main():
    x = np.array([0, 1, 2, ..., 9])
    y = np.array([1, 3, 2, ..., 12])

    b = estimate_coef(x, y)
    print(b)
```

→ Multiple

```
plot_r
```

Output :-

```
(b_0, b_
```

```
from sklea
    splot
import r
import n
from skle

data_wi

raw_dt

x = m

y = n

x_tra

reg =
reg

prin
```

plot_regression_line(x, y, b)

Output :-

$(b_0, b_1) = (1.2363\ldots, 1.16969\ldots)$

→ Multiple Linear Regression

```
from sklearn.model_selection import train_test
    split
import matplotlib.pyplot as plt
import numpy as np
from sklearn import datasets, linear_model,
    metrics

data_url = "http://lib.stat.cmu.edu/datasets/
    boston"
raw_df = pd.read_csv(data_url, sep="\s+",
    skiprows=22, header=None)

x = np.hstack([raw_df.values[::2, :],
    raw_df.values[1::2, :2]])
y = raw_df.values[1::2, 2]

x_train, x_test, y_train, y_test = train_test_
    split(x, y, test_size=0.4,
    random_state=1)

reg = linear_model.LinearRegression()
reg.fit(x_train, y_train)

print("Coefficient:", reg.coef_)
```

```
print("variance score: {}".format
    (reg.score(x_test, y_test)))

plt.style.use('fivethirtyeight')

plt.scatter(reg.predict(x_train), reg.predict
    (x_train) - y_train, color="green",
    s=10, label="Train data")

plt.scatter(reg.predict(x_test),
    reg.predict(x_test) - y_test,
    color="blue", s=10, label="test data")

plt.hlines(y=0, xmin=0, xmax=50, linewidth=2)

plt.legend(loc="upper right")

plt.title("Residual errors")

plt.show()
```

25/4/24

NP
25/4/24

25/4/24

Deci

01) Readi

02) Im

def

Aver

Inf

25/4/24

Date
Page

Decision tree ID3 implementation using play tennis

01) Reading data

02) Implementation

```
def find_entropy(df):
    target = df.keys()[-1]
    entropy = 0
    values = df[target].unique()
    for value in values:
        fraction = df[target].value_counts()[value] / len(df[target])
        entropy += -fraction * np.log2(fraction)
    return entropy
```

Average Information

$$I(\text{Attribute}) = \sum \frac{p_i + n_i}{p+n} \text{ Entropy(Attribute)}$$

Information gain Gain

Gain = Entropy(S) - I(Attribute)

Build Decision Tree

```
def buildTree(df, tree=None):
    target = df.keys()[-1]
    node = findwinner(df)
    attvalue = np.unique(df[node])
    if tree is None:
        tree = {}
        tree[node] = {}
    for value in attvalue:
        subtable = get_subtable(df, node, value)
    cvalue, counts = np.unique(subtable[target],
                               return_counts=True)
    if len(counts) == 1:
        tree[node][value] = cvalue[0]
    else:
        tree[node][value] = buildTree(subtable)

    return tree

Tree = buildTree(df)
```

Output :- `{outlook : {'overcast' : 'yes',`
`'Rain' : {'wind' : {'strong' : 'NO',`
`'weak' : 'Yes'}},`
`'Sunny' : {'humidity' : {'High' : 'NO',`
`'Normal' : 'Yes'}}}}`

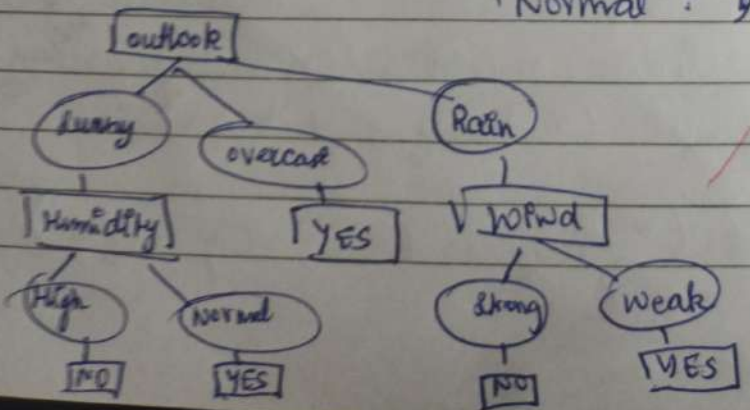

9/5/24

# Logistic Regression.

```
import pandas as pd
from matplotlib import pyplot as plt

file = "Insurance.csv"
df = pd.read_csv(file)

from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(x_train, y_train)
print(x_test)

y_predicted = model.predict(x_test)
print(y_predicted)
```

output :- Age : 22
probability : 0.1059

9/5/24

# KNN

```
import pandas as pd
from sklearn.datasets import load_iris
iris = load_iris()

df = pd.DataFrame(iris.data, columns=iris.feature_names)
df['target'] = iris.target
df['flower_name'] = df.target.apply(lambda x : iris.
                                    target_names[x])
```

```
from sklearn.model_selection import train_test_split
x = df.drop(['target', 'flower_name'], axis='columns')
y = df.target

x_train, x_test, y_train, y_test = train_test_split
                                   (x, y, test_size=0.2)

from sklearn.neighbors import KNeighboursClassifier
knn = KNeighborsClassifier (n_neighbors=10)
knn.fit(x_train, y_train)

knn.score(x_test, y_test)
// 0.9666

knn.predict([[4.8, 3.0, 1.5, 0.3]])
```

Output :- // 0 (setosa)

NLP 9/3/24

Lab 6- 07

K means Implementation

```
import pandas as pd
data = pd.read_csv("iris.csv")
data.head()

import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.cluster import KMeans

x, y = load_iris(return_x_y = True)

x means = KMeans(n_cluster = 3, random_state = 2)
KMeans.fit(x)

pred = kmeans.fit_predict(x)
pred
```

SVM

```
from sklearn.datasets import load_breast_cancer
import matplotlib.pyplot as plt
from sklearn.inspect import DecisionBoundaryDisplay
from sklearn.svm import SVC

cancer = load_breast_cancer()
x = cancer.data[:, :2]
y = cancer.target

svm = SVC(kernel="rbf", gamma=0.5, C=1.0)
svm.fit(x, y)

DecisionBoundaryDisplay.from_estimator(
    svm,
    x,
    response_method="predict",
    cmap = plt.cm.Spectral,
    alpha = 0.8,
    xlabel = cancer.feature_names[0],
    ylabel = cancer.feature_names[1],
)

plt.scatter(x[:,0], x[:,1], c=y, s=20,
            edgecolors="k")

plt.show()
```

# PCA

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
% matplotlib inline
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScale

from sklearn.datasets import load_breast_cancer
data = load_breast_cancer()
data.keys()

print(data['feature_names'])

df1 = pd.DataFrame(data['data'], columns=
                   data['feature'])

Scaling = StandardScalar()
Scaling.fit(df1)
Scaled_data = Scaling.transform(df1)
principal = PCA(n_components = 3)
principal.fit(Scaled_data)
X = principal.transform(Scaled_data)
```

30/8/21

30/5/24

Lab :- 08

→ Random Forest Ensemble method & Ada Boost

```
from sklearn.datasets import make_moons
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier,
    AdaBoostClassifier

x, y = make_moons(n_samples = 10000, noise = .5,
                  random_state = 0)

x_train, x_test, y_train, y_test = train_test_split
                  (x, y, test_size = 0.2,
                   random_state = 42)

clf = RandomForestClassifier(n_estimators = 100,
      max_features = "auto", random_state = 0)
clf.fit(x_train, y_train)

y_pred = clf.predict(x_test)

accuracy_score(y_test, y_pred)
```

Output :- 0.7965

30/5/24

Adaboost :-

```
clf. = AdaBoost Classifier (n_estimater =100)
clf. fit (x_train, y_train)
y_pred = clf. predict (x_test)

accuracy_score (y_test, y_pred)
```

Output :- 0.8333

Date 21/3/24
Page

Import and export pandas library functions

```
import pandas as pd

col_names = ["sepal_length_in_cm", "sepal_width_in_cm",
             "petal_length_in_cm", "petal_width_in_cm",
             "class"]

data = pd.read_csv("C:\\Users\\Admin\\Desktop\\
                    Iris\\Iris.data")

data.columns = col_names
```

→ output :-

| | sepal_length_in_cm | sepal_width_in_cm | petal length | petal width | Class |
|---|---|---|---|---|---|
| 0 | | | | | |
| 0 | 4.9 | 3.0 | 1.4 | 0.2 | Iris setosa |
| 1 | 4.7 | 3.2 | 1.3 | 0.2 | Iris-setosa |

21/3/24

Pandas
lessons

1. Creating, Reading & Writing
2. Indexing, Selecting & Assigning
3. Summary Functions & Maps
4. Grouping & Sorting
5. Data types & Missing Values
6. Renaming & Combining

Completed the Kaggle Pandas Certificate